

## 9. Rendimiento en Python

May 5, 2014

Dicen por ahí que Python es lento. Y lo es, a veces (en unas situaciones particulares): básicamente en bucles / ciclos

Un ejemplo:

```
In [4]: total = 0
        for i in xrange(100):
            total += i

        print total
```

```
File "<ipython-input-4-f736e62d08ed>", line 5
i = "hola
    ^
SyntaxError: EOL while scanning string literal
```

```
In [9]: total = 0
        total2 = ""
        for i in xrange(10):
            total += i
            print i

            i = "hola"
            total2 += i
            print i

        print total
        print total2
```

```
0
hola
1
hola
2
hola
3
hola
4
hola
5
hola
6
```

```

hola
7
hola
8
hola
9
hola
45
holaholaholaholaholaholaholaholaholahola

```

La flexibilidad misma lleva a la lentitud aquí, por ¡la necesidad de checar el tipo de los objetos todo el tiempo!

## 0.1 Diagnóstico por perfilamiento

Queremos acelerar un código de cómputo científico con muchos bucles. ¿Cuál es el procedimiento? Por ejemplo, “eliminar” los bucles. Antes que esto: veamos cuánto tarda en cada parte del programa. ¡Pero no a mano, sino con un *perfilador*!

La manera más fácil en IPython es con `%time` o `%timeit` (son diferentes): `%time` corre una vez un pedazo de código, mientras `%timeit` hace estadística. Y hay versiones con comandos mágicos de celda: `%%time` y `%%timeit` (corren todo el código de la celda completa):

```

In [10]: def sumar(n):
          total = 0
          for i in xrange(n):
              total += i

          return total

```

```

In [40]: %time sumar(10000000)

```

```

CPU times: user 624 ms, sys: 8.84 ms, total: 633 ms
Wall time: 626 ms

```

```

Out[40]: 49999995000000

```

```

In [41]: %timeit sumar(10**8)

```

```

1 loops, best of 3: 6.3 s per loop

```

```

In [20]: %%timeit
          total = 0
          for i in xrange(1000):
              total += i

```

```

10000 loops, best of 3: 57.4 µs per loop

```

Sin usar las bondades de IPython:

```

In [33]: import profile

```

```

In [34]: profile.run[?]

```

```

In [38]: profile.run("sumar(100000)")

```

4 function calls in 0.012 seconds

Ordered by: standard name

| ncalls | tottime | percall | cumtime | percall | filename:lineno(function)                |
|--------|---------|---------|---------|---------|--|
| 1      | 0.000   | 0.000   | 0.000   | 0.000   | :0(setprofile)                           |
| 1      | 0.012   | 0.012   | 0.012   | 0.012   | <ipython-input-10-6939f211451f>:1(sumar) |
| 1      | 0.000   | 0.000   | 0.012   | 0.012   | <string>:1(<module>)                     |
| 0      | 0.000   |         | 0.000   |         | profile:0(profiler)                      |
| 1      | 0.000   | 0.000   | 0.012   | 0.012   | profile:0(sumar(100000))                 |

## 0.2 Vectorizar

Reemplazar bucles explícitos de Python con operaciones vectoriales sobre `arrays` de `numpy`: `numpy` está diseñado tal que sus operaciones son rápidas – todos los bucles (implícitos) están escritos en C

```
In [42]: import numpy as np
```

```
In [49]: %%timeit
          #a = np.arange(10**6)
          s = sum(xrange(10**6))
          print s
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]

499999500000

100 loops, best of 3: 9.98 ms per loop

```
In [48]: %timeit sumar(10**6)
```

10 loops, best of 3: 62.1 ms per loop

```
In [53]: a = np.linspace(1, 10**6, 10**6)
         b = np.arange(0.1, 1, 10**6)
```

```
In [55]: %timeit c = a + b
```

1000 loops, best of 3: 1.24 ms per loop

```
In [59]: %%timeit
         c = np.zeros(10**6)

         for i in range(10**6):
             c[i] = a[i] + b[i]
```

-----  
IndexError

Traceback (most recent call last)

```
<ipython-input-59-6fa598d3b2ad> in <module>()
----> 1 get_ipython().run_cell_magic(u'timeit', u'', u'\nc = np.zeros(10**6)\n\nfor i in range(10**6)
```

```

/Users/dsanders/development/ipython/IPython/core/interactiveshell.pyc in run_cell_magic(self, magic_name, line, cell)
2136         magic_arg_s = self.var_expand(line, stack_depth)
2137         with self.builtin_trap:
-> 2138             result = fn(magic_arg_s, cell)
2139         return result
2140
```

```

/Users/dsanders/development/ipython/IPython/core/magics/execution.pyc in timeit(self, line, cell)
```

```

/Users/dsanders/development/ipython/IPython/core/magic.pyc in <lambda>(f, *a, **k)
189     # but it's overkill for just that one bit of state.
190     def magic_deco(arg):
--> 191         call = lambda f, *a, **k: f(*a, **k)
192
193         if callable(arg):
```

```

/Users/dsanders/development/ipython/IPython/core/magics/execution.pyc in timeit(self, line, cell)
984         number = 1
985         for i in range(1, 10):
--> 986             if timer.timeit(number) >= 0.2:
987                 break
988             number *= 10
```

```

/usr/local/Cellar/python/2.7.5/Frameworks/Python.framework/Versions/2.7/lib/python2.7/timeit.py
```

```

193         gc.disable()
194         try:
--> 195             timing = self.inner(it, self.timer)
196             finally:
197                 if gcold:

```

```

<magic-timeit> in inner(_it, _timer)

```

```

IndexError: index 1 is out of bounds for axis 0 with size 1

```

```

In [57]: len(a)

```

```

Out[57]: 1000000

```

```

In [63]: a = 3
        s = "print(a)"

```

```

In [64]: s

```

```

Out[64]: 'print(a)'

```

```

In [66]: exec(s)

```

```

3

```

```

In [67]: expr = "3*a"
        eval(expr)

```

```

Out[67]: 9

```

### 0.3 List comprehensions

Supongamos que queremos aplicar una función a cada elemento de una lista, por ejemplo duplicar cada elemento:

```

In [68]: a = np.arange(10)
        b = 2*a

```

```

In [69]: a

```

```

Out[69]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

```

```

In [70]: b

```

```

Out[70]: array([ 0,  2,  4,  6,  8, 10, 12, 14, 16, 18])

```

```

In [71]: def f(x):
        return x*x + 1

```

```

In [72]: map(f, a)

```

```

Out[72]: [1, 2, 5, 10, 17, 26, 37, 50, 65, 82]

```

```

In [73]: f(a)

```

```

Out[73]: array([ 1,  2,  5, 10, 17, 26, 37, 50, 65, 82])

```

List comprehension:

Matemáticamente, escribiríamos:

$$\{x^2 + 1 : x \in A\}$$

```
In [74]: [x*x+1 for x in range(10)]
```

```
Out[74]: [1, 2, 5, 10, 17, 26, 37, 50, 65, 82]
```

```
In [92]: %time l = [x*x+1 for x in xrange(1000000) if x%2==1]
```

CPU times: user 223 ms, sys: 11.7 ms, total: 235 ms

Wall time: 229 ms

```
In [95]: %%timeit
l = []
for x in xrange(1000000):
    if x%2==1:
        l.append(x*x+1)
```

10 loops, best of 3: 165 ms per loop

## 0.4

```
In [78]: np.sin(a)
```

```
Out[78]: array([ 0.          ,  0.84147098,  0.90929743,  0.14112001, -0.7568025 ,
                -0.95892427, -0.2794155 ,  0.6569866 ,  0.98935825,  0.41211849])
```

```
In []:
```

## 0.5 Siguiente paso: Cython y numba

Cython y numba compilan pedazos de Python a C

```
In []:
```